Page-event admissions enrichment: the visitor sample

Story `admissions.CAP.page_event_admissions_enrichment` · Done means: "20 sampled visitor summaries match the underlying page events with zero fabricated page titles."

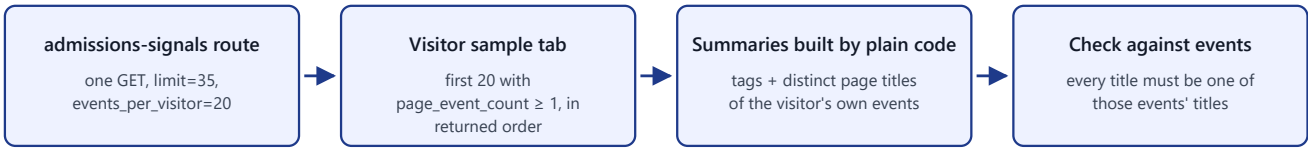
● Criterion met ● Read-only, no writes

20 of 20 match · 0 fabricated titles

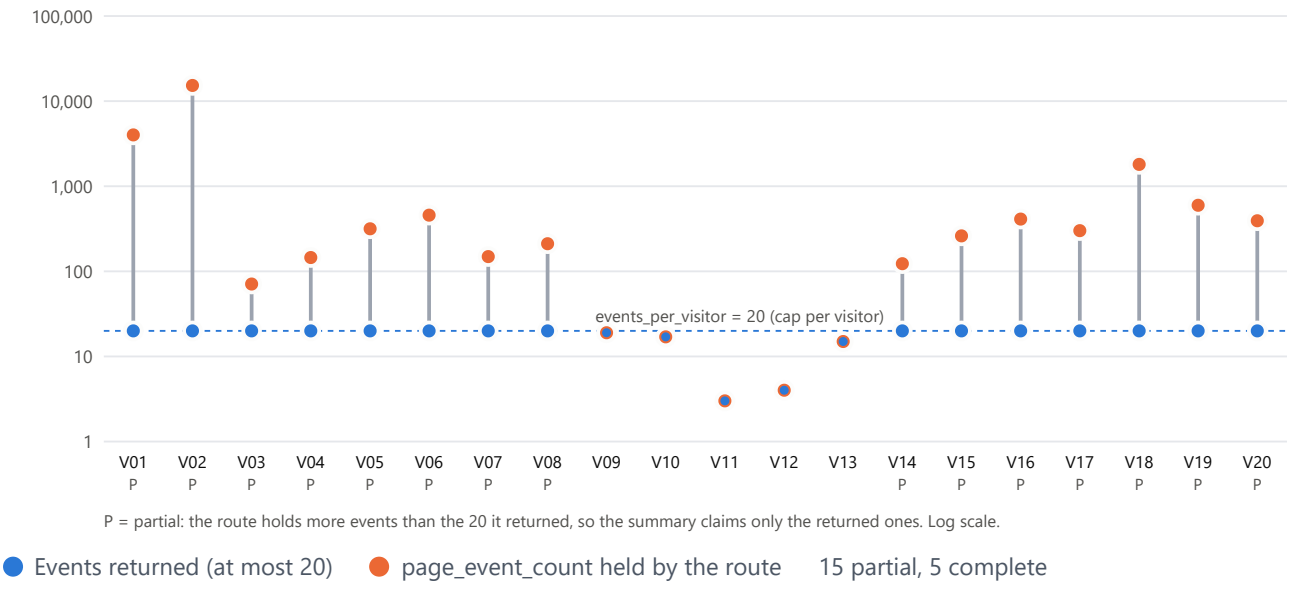
Pulled 12 November 2026, 2:39 PM CST (HTTP 200) `limit=35` `events_per_visitor=20`

35	20	12	3
visitors returned by the route	sampled: the first 20 with page_event_count ≥ 1	skipped for no page events	not reached after the sample filled

How a summary is made and checked



Events returned vs page_event_count, per visitor



Where the two dots meet, every event the visitor has was returned. Where a grey line joins them, the visitor is partial and the summary claims only the 20 returned events. All 20 rows: Match, 0 fabricated titles.

The tab, as it showed the pull

Visitor sample

Evidence for admissions.CAP.page_event_admissions_enrichment. Done means: "20 sampled visitor summaries match the underlying page events with zero fabricated page titles."

Read-only. Summaries are built and checked by plain code, not by AI. One pull of GET /api/admin/agent-domains-mine/admissions/admissions-signals?limit=35&events_per_visitor=20. From that single response, the visitors whose page_event_count is at least 1, in the order the route returned them, up to 20. Visitors with no events are skipped and counted. If fewer than 20 have events, the sample is reported as it is: nothing is pulled again, and never with a limit above 35 or events_per_visitor above 20. A visitor with more events than the 20 returned is marked partial.

Signed in. The pull uses the admin token minted on the Agent domain tab, held in this page's memory only.

`GET /api/admin/agent-domains-mine/admissions/admissions-signals?limit=35&events_per_visitor=20`

Pulled

20 of 20 match · 0 fabricated titles

Pulled at 12 November 2026, 2:39 PM CST (HTTP 200). One pull per page load: this sample is fixed and is not pulled again.

- Visitors in the response: 35 (the route's "returned": 35; "matching", its count of admissions memory rows: 116).
- Skipped for having no page events: 12.
- Skipped for a missing field: 0.
- Not reached after the sample filled: 3.

What the route says about itself, as returned

> Anonymised by design. anon_id is a keyed HMAC of the visitor id, stable within this surface and not reversible to a visitor; it rotates if the signing secret rotates, so nothing should be joined on it across a rotation. page_path and page_title are returned in place of page_url, which carries the query string the tracker writes an email into. page_title is null when the event recorded none, never an empty string, because the capability this serves is measured on fabricating no titles.

not_recorded: ip_address, user_agent, fingerprint, lead_id, page_url, event_data, questions_asked, conversation_summaries, personality_notes

The Visitor sample tab after the pull: the totals line, the counts, and what the route says about itself.

Results

Visitor	Events returned	page_event_count	Partial	Titles in summary	Fabricated titles	Result
V01	20	4008	partial	1	0	Match
V02	20	15283	partial	1	0	Match
V03	20	71	partial	1	0	Match
V04	20	145	partial	1	0	Match
V05	20	317	partial	1	0	Match
V06	20	458	partial	1	0	Match
V07	20	149	partial	1	0	Match
V08	20	211	partial	2	0	Match
V09	19	19		1	0	Match
V10	17	17		1	0	Match
V11	3	3		1	0	Match
V12	4	4		1	0	Match
V13	15	15		1	0	Match
V14	20	123	partial	1	0	Match
V15	20	261	partial	1	0	Match
V16	20	411	partial	1	0	Match
V17	20	300	partial	1	0	Match
V18	20	1804	partial	1	0	Match
V19	20	598	partial	1	0	Match
V20	20	392	partial	1	0	Match

Download results

The results table, V01–V20. No ids, page paths, page titles or tags are shown.

Route returns / never returns

The route's own disclosure and `not_recorded` list, as it returned them in this pull.

Returns

`anon_id`, a keyed HMAC
Stable within this surface and not reversible to a visitor; rotates if the signing secret rotates.

`page_path` and `page_title`
In place of `page_url`, which carries the query string the tracker writes an email into.

`page_title` is null when none was recorded
Never an empty string. The tab shows it as "(no title)".

Never returns (`not_recorded`)

- `ip_address`
- `user_agent`
- `fingerprint`
- `lead_id`
- `page_url`
- `event_data`
- `questions_asked`
- `conversation_summaries`
- `personality_notes`

> "Anonymised by design. anon_id is a keyed HMAC of the visitor id, stable within this surface and not reversible to a visitor; it rotates if the signing secret rotates, so nothing should be joined on it across a rotation. page_path and page_title are returned in place of page_url, which carries the query string the tracker writes an email into. page_title is null when the event recorded none, never an empty string, because the capability this serves is measured on fabricating no titles."

How it was kept honest

One pull only
One GET per page load. A finished pull is reused, and it is never pulled again with a bigger limit.

No joining on anon_id
Each visitor is matched only to the events nested inside it in the same response; anon_id is never a join key, within a pull or across pulls.

Missing titles shown as "(no title)"
A null page_title is never filled in from page_path or anything else.

Summaries built without AI
Plain code builds each summary from the visitor's own events and checks it separately against those events.

The 17 Visitor sample tests, all passing (`node --test cc/tests/visitorSample.test.mjs`):

- the route path is the one the platform merged in PR 3779, with limit=35 and events_per_visitor=20
- only limit=35 and events_per_visitor=20 are allowed: never a limit above 35 or events_per_visitor above 20
- 35 returned: the first 20 with events, in returned order, are sampled and the rest are not reached
- fewer than 20 of the 35 returned have events: reported plainly, and not pulled again with a bigger limit
- a matching summary passes
- an invented title is caught as fabricated
- a missing page_title shows "(no title)" and is never filled in from page_path
- visitors with zero events are skipped and counted, and the first 20 in order are kept
- fewer than 20 visitors with events is reported as short, not padded
- a visitor with more events than were returned is marked partial and claims only those
- no second pull happens: a finished pull is reused and concurrent clicks share one call
- a failed pull samples nothing and shows the error, and may be tried again
- visitors are never joined across pulls, or by anon_id within one
- a missing field is shown by name and nothing is filled in
- the download holds the table, totals, pull time, limits and skipped count, and no ids, paths, titles or tags
- the pull uses the minted admin token once, with no retry, and never shows it
- the pull is refused before sign-in, without calling the portal

Found along the way

The 11 November, 8:14 PM CST failure was a deploy restart
Confirmed by the platform team.

The unbounded page-event query is partly fixed
Reported from this work. The platform team added an index and a per-visitor cap, both live, so the pull is safe. The query still sorts every matching row before the cut; the platform team is rewriting it to fetch per visitor.

`limit=35` was measured and approved
By the platform team, with `events_per_visitor=20`. The tab refuses any other value.

Open items (platform team's tickets)

The agent doesn't match the capability brief
The agent doesn't behave as the brief describes: the 15-minute read has no watermark, there is no 'unclassified' result, and an error on the activity log write is swallowed. No test covers it. The platform team holds the ticket.

The R8 citation belongs to `agentDomainReleaseLadder.ts`
Not to the file the brief cites; the platform team holds the ticket.

Built only from `sample-results.json` and the two screenshots of the Visitor sample tab. No anonymous ids, page paths, page titles, tags or personal data. The results file and screenshots are not committed.